

Consent and the Right to Privacy

KEVIN MILLS

ABSTRACT *There is currently intense debate about the significance of user consent to data practices. Consent is often taken to legitimate virtually any data practice, no matter how invasive. Many scholars argue, however, that user consent is typically so defective as to be 'meaningless' and that user privacy should thus be protected by substantive legislation that does not rely (or does not rely heavily) on consent. I argue that both views rest on serious mistakes about the validity conditions for consent. User consent is sufficiently impoverished that it does not guarantee legitimacy but is not so impoverished as to be 'meaningless'; it can legitimate data practices that are independently reasonable but not those that are exploitative. Since many valuable data practices must be consented to if they are to be legitimate (or so I argue), our privacy legislation should continue emphasizing the importance of user consent, even if auxiliary protections are also desirable.*

Almost everything people do online is tracked by a variety of entities, whether they are shopping, using a search engine, reading the news, getting health information, or even watching pornography. One regulatory model for protecting people's privacy from all this tracking, and currently the dominant legal framework in the United States, is known as 'notice-and-consent'. Under notice-and-consent, people are mostly free to collect and use data as they wish, but they are required to notify users of their data policies, typically by including a link to a privacy policy. Users are then taken to consent to these, either expressly (by clicking 'I agree' or some cognate) or tacitly (by merely using the services).

According to advocates, notice-and-consent provides all the privacy protections people need; notice empowers users to decide for themselves whether data practices are reasonable, and their consent legitimates these practices. Critics have objected on several grounds. Daniel Susser offers a nice summary:

Philosophers and legal theorists argue that the cost of opting out of services is often too high to represent a meaningful choice for individuals to make. In cases where they might reasonably opt out individuals are frequently ill-equipped to fully understand the decision they face. And even when they do understand it, what is at issue in the decision is often a social rather than individual interest. In such cases individuals should not be deciding about it in the first place.¹

The upshot of these criticisms, Susser proposes, is that 'the consent offered when individuals sign privacy policies ... is not meaningful' and 'cannot legitimate the collection and use of user information'.² Susser thus joins a growing chorus of scholars who propose that consent to data practices is 'meaningless'.³ Importantly, critics typically claim that these problems arise from structural and technical features of data practices that are either

impossible or extremely unlikely to change, and that as a result, 'Notice-and-consent appears to be a failed regulatory model'.⁴

Against all this, I argue that advocates and critics of notice-and-consent are both mistaken about the significance of consent to data practices. Advocates are wrong to think that consent guarantees legitimacy, for reasons critics have rightly emphasized. But critics are wrong to think that consent 'is not meaningful' and 'cannot legitimate' data practices; while consent does not guarantee legitimacy, it may legitimate all the same. The basic reason is that if notice were improved, consent would neither be so impoverished as to be meaningless nor so impeccable as to guarantee legitimacy; it could legitimate data practices that are independently reasonable but not those that are exploitative.

The major upshot is that people often can provide meaningful consent to data practices – even complicated data practices that are costly to avoid. This does not mean that notice-and-consent is a good regulatory model; even if people can provide meaningful consent to data practices, it may often be preferable to spare them the burdens of doing so (and/or to strengthen their bargaining positions by denying them the ability to do so).⁵ But it does mean that invasive data practices (i.e. data practices which would conflict with an individual's rights) can, in the right circumstances, be legitimated by consent. This in turn means that where there are (on balance) genuinely good reasons to deploy such data practices, this can be done in good conscience, as long as consent is properly secured and care is taken to not exploit data subjects.

1. Framing the Debate

The function of consent is to voluntarily alter moral and/or legal (henceforth: 'normative') entitlements. For our purposes, this chiefly involves waiving some of our rights (thus releasing certain others from any correlative duties) and/or undertaking obligations (either to the performance or nonperformance of certain actions). Our ability to alter normative entitlements is, of course, not unlimited; I typically cannot waive your rights, for example. Where somebody's consent can alter some normative entitlement (when expressed through suitable means, which can vary between contexts), I will say that they have a normative power. One's normative powers thus demarcate a space within which one can control rights and duties, sometimes one's own, sometimes other people's.⁶ Call consent that does (morally) or should (legally) alter entitlements in the ways it purports to 'valid'.

In the present context, the relevant function of consent is rendering permissible certain data practices which would otherwise be impermissible. Structurally, this is best understood as involving two components. First, data practitioners (people who engage in data practices) initially have duties, owed to technology users, to refrain from certain data practices (equivalently: users have rights, held against data practitioners, to be free from certain data practices).⁷ Second, users have powers to waive these duties, thus permitting (*ceteris paribus*) data practitioners to do something they would otherwise not be permitted to do. It is only when (and because) there is this distribution of duties and powers that user consent to data practices is significant; where there are no initial duties to refrain from data practices, user consent is unnecessary, and where users lack the powers to waive such duties, user consent is irrelevant.

It will be helpful to have a label in what follows, so I will call any duty prohibiting a data practice a 'data restriction'. The question of the proper role for user consent to data

practices is fundamentally about whether and when users should be able to waive data restrictions (versus when, perhaps for the protection of users themselves, certain data restrictions should be unwaivable). This breaks down into two sub-questions:

POWERS: Which data restrictions should users have the legal powers to waive?

CONDITIONS: When should users be recognized as exercising these powers?

It is worth emphasizing that *CONDITIONS* is distinct from *POWERS*; there is sometimes a temptation to speak only of 'consent' in the abstract, but we can, should, and do adopt different standards for the conditions under which different powers are exercised (for example, legally binding contracts have various formal requirements, and we have uncommonly strict requirements for consent to medical research).

Much of the theoretical work on consent to data practices has been done in the context of criticizing notice-and-consent, which offers particularly promiscuous answers to these questions; notice-and-consent grants users nearly unlimited powers to legitimate data practices and places almost no restrictions on the conditions under which these powers can be exercised. Two standard criticisms have emerged. The first claims that user consent is typically 'meaningless', as users are uninformed (they cannot reasonably be expected to understand what they are consenting to) and their consent is involuntary (they typically have 'no choice' but to consent even if they do). Structurally, this challenges notice-and-consent's answer to *CONDITIONS*. That said, critics typically argue that these problems cannot realistically be addressed and use this to challenge notice-and-consent's answer to *POWERS*, arguing for a variety of data restrictions users are not empowered to waive.

The technical details of these arguments are typically made much more explicit than the normative ones. To establish that user consent is involuntary, critics note that invasive data practices are so pervasively tied to important technologies (e.g. the internet) that avoiding them essentially requires forgoing these technologies altogether, something thought so unrealistic as to invalidate any consent offered to bundled data practices.⁸ To establish that users are uninformed, critics cite a variety of features that allegedly make 'informed consent' to data practices difficult or impossible, including: (i) the length and complexity of privacy agreements, both as they actually are, and as they allegedly must be in order to adequately inform users;⁹ (ii) the sheer volume of privacy agreements users encounter;¹⁰ (iii) the ability of data practitioners to unilaterally update privacy agreements at any time, thus requiring constant monitoring from users;¹¹ (iv) the fact that the inferences that can be drawn from data, and thus the costs of giving up this data, can be highly counterintuitive;¹² (v) the frequency with which data flows to third parties (e.g. is sold on data markets), which often makes it all but impossible to assess who will ultimately gain access to one's data and what they will do with it;¹³ and (vi) the aggregative nature of data costs (new data can combine with old data in complicated ways to facilitate even more counterintuitive inferences).¹⁴ All told, or so the argument goes, assessing the costs of data practices is so difficult that most users cannot offer 'meaningful' consent to them.

We can summarize the first criticism as arguing that user consent is defective. The second criticism argues that user consent is, in many cases, irrelevant. In granting users nearly unlimited powers to legitimate data practices, notice-and-consent presupposes that data restrictions are owed almost exclusively to individual users, but critics reject this. The standard argument is that many data restrictions are justified by broad social

considerations (for example: because data is being used to manipulate people in ways that undermine democracy, and this must be stopped). That said, we need not appeal to broad social values to make this point; if there is any justification for data restrictions other than the protection of the individual user whose data is at issue (including e.g. the protection of groups smaller than society generally, such as marginalized groups who may be discriminated against), notice-and-consent's answer to *POWERS* fails, since users should not have powers to waive duties that are not owed to them.

While these criticisms are surely on to something, I believe they are less forceful than is standardly believed; they rightly highlight various respects in which notice-and-consent may leave people vulnerable to unscrupulous data practitioners, but they are unreasonably pessimistic about what consent can accomplish when data practices are scrupulously deployed in good faith (including, perhaps, when doing so is enforced by market conditions rather than legislation). Let us consider each criticism in more detail.

2. Information

Let us start with the first criticism, which argues that consent under notice-and-consent is typically invalid since: (i) users cannot reasonably be expected to understand what they are consenting to; and (ii) they typically have 'no choice' but to consent even if they do. As Alan Wertheimer's work in particular has demonstrated, these arguments are considerably more complicated than is sometimes appreciated.¹⁵ I will start with information.

Conventional wisdom has it that valid consent either should or must be 'informed'. It is not always clear exactly what this means, but on a fairly standard analysis, it requires knowing all the following (to some suitable level of depth):¹⁶

- (1) what is being consented to;
- (2) the likely costs and benefits of consenting;
- (3) the details of any relevant alternative courses of action.

The standard argument for this sort of view appeals to autonomy (or to a close legal correlate, 'the right to self-determination').¹⁷ For example, Faden and Beauchamp argue that autonomy ought to serve as 'the benchmark' by which consent is evaluated and that 'An action cannot be autonomous if the actor fails to have an understanding of his or her action'.¹⁸ Similarly, Barocas and Nissenbaum propose that 'There is little value in a protocol for informed consent that does not meaningfully model choice and, in turn, autonomy', something that requires people to 'understand how their assent plays out in terms of specific commitments, beliefs, needs, goals, and desires'.¹⁹ The core idea seems to involve three claims: first, that the purpose of soliciting consent (the only purpose?) is to facilitate autonomous decision-making; second, that autonomous decisions must be appropriately responsive to an agent's values; and third, that this requires the agent to be informed in roughly the previous sense.

There is surely much to recommend this proposal, but it belies a great deal of complexity. To start, we should distinguish two possible claims: first, that valid consent must be informed, and anything insufficiently informed is thus, by that fact alone, not valid consent; second, that valid consent should ideally be informed, but in ways that might be outweighed by other considerations. Let us consider each option in turn.

To help explore the proposal that valid consent must be informed, consider the following example:

CULPABLE IGNORANCE: Culpable wants to rent an apartment and is given a lease. She could easily read and understand the lease, but chooses not to, because she prefers to watch television. Culpable signs and returns the lease the next day, falsely claiming that she has read it and agrees to the terms. Landlord stops showing the apartment to other prospective tenants, believing a tenancy agreement has been formed.

Culpable is significantly uninformed about the lease she has signed but has nobody to blame but herself. Assuming the terms are reasonable, she should be held to them, even if she did not know about them, and even if she would not have signed if she had. For example, imagine the lease prohibits pets, whereas Culpable was planning to purchase a cat. It is unfortunate that she is now bound by this prohibition, but that is her burden to bear; to treat her consent as invalid would be unfair to Landlord.²⁰

To prevent misunderstandings, the lesson of *CULPABLE IGNORANCE* is not that users should be held responsible for failing to understand privacy agreements (a naïve proposal at best, a malicious one at worst).²¹ The lesson is that if there is some sense in which valid consent must be informed, this involves informational requirements far lower than those standardly associated with 'informed consent'. Culpable does not know exactly which entitlements her consent is altering. If you add my addendum about pets, she also does not know about an important cost of consenting. And while this was not part of the example, she could easily know very little about the available alternatives, e.g. because this was the only apartment she looked into. None of this invalidates her consent.

In this example, the most important reason to treat Culpable's consent as valid is that this is a bilateral consent transaction, and fairness to the other party requires doing so. It is worth emphasizing that this is not the only sort of case. Consider another example, this time involving marriage (which transforms normative entitlements in various important ways, e.g. pertaining to property, decision-making powers in the event of incapacitation, the fiduciary obligations that come with these, etc.)²²:

OPTIMISTIC MARRIAGE: Optimist and Idealist are in love and decide to get married. They know that many marriages end in divorce; however, they are all but certain this will not happen to them, as their love is true love.

While some people in this (probably not uncommon) situation might be justified in their confidence, many presumably would not be. Assume Optimist and Idealist are not – assume their confidence is not premised on some deep insight into their relationship but is simply a by-product of their intoxicating love. This makes their decision to get married not just uninformed, but significantly misinformed, as they have false beliefs about the probabilities of various costs and benefits the decision implicates. Even still, I propose, their consent is and should be treated as valid; we should neither prevent them from marrying, nor, if their marriage fails, retroactively declare their marriage null and void on the grounds that it resulted from faulty consent.

To be clear, I do not take any of this to refute the proposal that valid consent must be informed in some important but more minimal sense. Grounds for doubt notwithstanding,²³ I find this extremely plausible, for reasons Tom Dougherty nicely articulates: consent must typically be intentional, and one must have at least some

understanding of what one is doing for an action to be intentional.²⁴ What is important for my purposes is that if this is true at all, these requirements fall far short of the standard requirements for ‘informed consent’; they may include various details about what is consented to, but will not, at least in general, include a deep understanding of costs and benefits or available alternatives.

I turn now to the second proposal: that valid consent should ideally be informed, but only as one consideration amongst others. Perhaps this is correct – either because of some deep truth about autonomy and consent, or simply because being informed makes people more likely to exercise their powers in ways that promote their welfare. The important thing to recognize is that there are considerations that pull in the opposite direction. We are limited beings living in complex environments, and there are many reasons why we may not wish or be able to thoroughly inform ourselves about various consent decisions. When such circumstances arise, strict information requirements can deny us the power of choice entirely, or at least make this power very difficult to exercise. While this is surely sometimes appropriate, it often has serious costs; denying people the power of choice either means denying them certain opportunities, which is problematic when these opportunities are valuable, or forcing these opportunities upon people, which is problematic when their desirability or legitimacy depends on people’s consent. Given these costs, it is sometimes better to grant people the power of choice even where a decision involves serious information deficits.

Consider marriage again. Marriage is a serious legal arrangement with serious legal repercussions, yet decisions to get married often involve significant information deficits; the costs and benefits are uncertain, the probabilities are difficult to assess, the alternatives are massively unknown and poorly explored, and the decision implicates powerful emotions that are known to distort our judgment. Why do we let people enter into such a serious legal arrangement with such impoverished information (and not just let them, but even incentivize doing so with various legal benefits)? I propose that the best justification at least includes the following: because we think marriage is an important and valuable institution, for the married couple and/or for society generally. If it were not (or, I suppose, if you think it is not), it might well be preferable to deny people the power to marry on grounds of protecting them from bad decisions resulting from weak information.

For another example, consider medical research, which inevitably involves significant information deficits and sometimes subjects participants to serious risks. It is hard to justify allowing participants to consent to this except by appealing to the value of medical research, and indeed, our regulatory structures presuppose this (institutional review boards are tasked with assessing whether a study’s potential benefits justify its risks to participants and will typically prohibit the recruitment of human subjects for risky studies that lack ‘social value and scientific validity’).²⁵ Importantly, the value facilitated by medical research is not, or is not only, to the benefit of participants; one of the crucial justifications for allowing people to consent to risky studies despite serious information deficits is that this facilitates something of great importance to society generally. And, or so I would urge, this justification succeeds. As these information deficits cannot be eliminated (although we do the best we can), and as it would (I assume) typically be illegitimate to force people to participate in medical research, the value of medical research justifies lower information requirements for consent than might apply in other contexts.

My proposal, then (Wertheimer’s proposal, really), is that there are no context-insensitive information requirements for valid consent, except perhaps of a quite minimal

sort; the threshold of information beneath which people should be denied some power, or prevented from exercising it, depends on factors that are specific to the power in question. One factor is the comparative importance of protecting people from bad decisions versus facilitating good ones, something that depends on the stakes of the decision. Another pertains to the viability of bypassing consent altogether, i.e. whether it is possible, legitimate, and desirable to deny people some power and promote the values consent might otherwise facilitate through other means. This is impossible for values that constitutively depend on consent. It is illegitimate where it violates people's (moral) rights, either to have certain legal powers or to be free from certain non-consensual interactions. And it is undesirable where the values at stake would be more effectively promoted consensually, whether because people prefer to make decisions themselves, because their decisions track their preferences better than alternatives, or for some other reason.²⁶

Consider now consent to data practices. I grant that there are various technical and structural features of data practices that significantly limit people's abilities to assess the costs of consenting to them. If we add the plausible assumption that these costs can be significant (if only because people value their privacy, although perhaps also in more tangible ways), this makes a strong case for declaring such consent invalid on grounds of protection. But this is, as argued, only one part of the story; we must also consider the values consent to data practices potentially facilitates and the viability of bypassing consent altogether. I will consider these in reverse order.

Note first that users can be informed about data practices in various important ways (or at least presumably could be, if data practitioners honestly and accessibly summarized them). Most notably, they can often understand what data is being collected, certain salient features about how it is used, and who initially collects it (or embeds services that allow others to collect it). Each of these is extremely significant – they are some of the most significant features of many data practices. For example, many in-home assistants extensively record audio after an activation command is spoken (or is mistakenly flagged as spoken), and workers are sometimes employed to listen to this audio to improve voice recognition algorithms.²⁷ Many people find this extremely invasive, and not because of complex technical details they cannot be expected to understand; the recording itself is invasive, as is allowing people to listen to it. That said, there are good reasons for this practice (improving voice recognition algorithms is important), and many people seem untroubled by it, or at least willing to tolerate it in exchange for the benefits offered by the technology. Given this, I see no reason (at least based on concerns about information) why consent cannot legitimate these specific practices, following suitable disclosures. Moreover, I would urge that consent both should and must be solicited here; should, because it will help track people's preferences about whether they are subjected to these practices, and must, because it simply strikes me as illegitimate to extensively record audio inside of people's homes without their consent (I cannot develop a proper argument for this claim here, but I hope your good moral judgment suffices; this is clearly a practice that otherwise conflicts with the right to privacy, and that must thus be consented to if it is to be legitimate).

This is a fairly extreme case, but it is not a unique case. Do I want to use a search engine that indefinitely retains my search queries? Or map software that indefinitely retains my location data? Or a website with analytics software that produces (or can produce) reconstructed videos of my web browsing? Or cloud storage that searches through the contents of my photographs? Or an email provider that builds a profile on me by searching through

my emails? Or a social media platform, which extensively collects and uses data, run by a company with a dubious ethical history? I could go on, but I hope the point is already sufficiently clear: people can track many salient features, both of data practices and of those who engage in them, and often have strong and reasonable preferences because of this. At least if we restrict our attention to information, their consent to such practices is not at all 'meaningless'; people know enough, or can know enough, to bring meaningful and important preferences to bear on whether to use such technologies. Moreover, I would urge again that many of these practices must be consented to if they are to be legitimate.

I thus propose that it would be undesirable, and often illegitimate, to permit these data practices without user consent. Should they be permitted at all? There is a strong case to be made that many of them should; data is extremely valuable, and users derive both technological and economic benefits from data practices (both as individuals and as members of our society). Perhaps most obviously, data is one of the twin economic pillars of the internet that make so much content available for 'free' (the other is advertising, although it is somewhat misleading to separate these). To offer an example I find particularly striking, one can find instructional videos online, many of them quite high quality, covering almost any topic one might reasonably want instruction on (e.g. yoga, cooking, fixing your car, assembling your own computer, or even coaching Little League baseball). This knowledge is made available to everybody with an internet connection and is significantly monetized – both in distribution and often in production – by data practices. One argument in favor of permitting these data practices is that there is value in being able to pay for content with data instead of money, but it goes deeper than this. For one thing, it is at least unclear whether content would flourish to the extent it currently does under alternative economic models. For another, there is an important social-justice function served here, as the egalitarian distribution of specialized knowledge is no small matter. And this is to say nothing of the respects in which our technology itself is improved by data practices.

I am going to take a step back, as I cannot properly explore the considerations that count for and against various data practices in this article (to say nothing of the viability of alternative economic models). My fundamental proposal, however, is that this is the task that faces us. People have rights, most notably the right to privacy, and data practices sometimes conflict with these rights. Where they so conflict, these practices must be consented to if they are to be legitimate (the rights-holders must waive the rights in question). If people could not, in principle, offer valid consent here, they would either have to be denied certain rights, various data practices would inevitably be illegitimate, or both. But at least from the perspective of information, this need not be the case: first, because users can be informed about some very important features of data practices; second, because the requirements for valid consent depend on the comparative importance of protecting people from bad decisions versus facilitating good ones (whether good for the consentor, society generally, or both); and third, because invasive data practices sometimes have considerable value. Where there are good reasons to deploy certain data practices (economic, technological, or otherwise), where these reasons trump or outweigh countervailing reasons, and where sincere efforts are made to inform data subjects in whatever relevant ways are feasible, people can thus offer valid consent to data practices despite significant information deficits.

Of course, we must be careful here; insofar as people have weak information, they are poorly positioned to safeguard their interests, and vulnerable to exploitation. If we are

to accept their consent, it is morally imperative that their interests are somehow protected elsewhere – either by scrupulous data practitioners, a well-ordered market, regulators, or some combination of the above. But as long as they are properly protected, consent can legitimize practices that are otherwise illegitimate, and such practices can thus be deployed in good conscience (at least from the perspective of information). Whether and when data practices satisfy these requirements depends on difficult questions about their empirical effects, the right to privacy, distributive justice, and other matters. As this article is primarily on consent, I must set such issues aside.

3. ‘No Choice’ Scenarios

Let us now explore the claim that consent under notice-and-consent is often invalid because users have ‘no choice’ but to provide it. I will note in passing that the meaning of this expression is not transparent, and what appears to be the same basic idea is often expressed in different ways, e.g. by saying that somebody ‘lacks reasonable alternatives’, ‘lacks acceptable alternatives’, or ‘lacks meaningful choices’. That said, rather than attempt to clarify this idea, I will simply assume we have an adequate intuitive grasp on it, since I will argue that it is the wrong way to approach these matters anyway.

Scholars working on consent have long recognized that lacking acceptable alternatives, whatever exactly this means, does not in general invalidate somebody’s consent.²⁸ Consider a simple example:

SURGERY: Patient needs surgery and will otherwise die. The surgery is safe, effective, and relatively pain free. Patient consents to the surgery and Surgeon performs it.²⁹

If anybody lacks acceptable alternatives, Patient surely does – his only alternative is death. But this does not invalidate his consent, as his consent renders (or could render) the surgery permissible. What follows if he does not consent is complicated. Given my stipulation that the surgery is ‘safe, effective, and relatively pain free’, his refusal would arguably be evidence of incompetence or misinformation, and would leave Surgeon with some investigative work to do.³⁰ But these details do not matter for our purposes; after disclosure, his affirmative consent renders the surgery permissible (with no further investigations needed), and a lack of acceptable alternatives thus does not invalidate his consent.

There are two interesting structural features of *SURGERY*: first, consenting presumably is an acceptable alternative; second, the consent receiver is doing all they can to help the consent provider. It is worth considering an example that lacks these features. Consider:

RING: Seller needs money to pay for an urgent medical treatment. She believes her only realistic way to get this money is to sell her late husband’s wedding ring, which is quite valuable, but also of deep emotional significance. Seller locates Buyer, receives fair market value for the ring, and is thus able to afford the treatment.

Both options available to Seller are presumably unacceptable. Buyer also does less to help than is possible (Buyer could give Seller the money outright). Does this render Seller’s consent to sell her *RING* invalid? If it does, then Buyer violates Seller’s rights, as he takes

Seller's property without her valid consent. But Buyer surely does not violate Seller's rights here; he might be subject to other forms of criticism, but not this one. It follows that Seller's consent is valid, even though she lacks acceptable alternatives and even though Buyer could but does not offer her one.

I believe people are sometimes misled about the significance of lacking acceptable alternatives because situations in which somebody lacks acceptable alternatives often (but not always) have a separate feature which really does invalidate consent. Consider:

ROBBERY: Robber points a gun at Victim and says, 'your money or your life'.
Victim hands over his money.

Victim lacks acceptable alternatives, and his consent is surely invalid. But notice another feature of this example: Robber violates Victim's rights by demanding his compliance under the threat of murder, thus unjustly limiting the options Victim has available to him (Victim should have a third option where he keeps his money and lives). Following Wertheimer's analysis, I believe this rights violation is what undermines Victim's consent, not a lack of acceptable alternatives.³¹ Consider another example:

COFFEE: Bully grabs Innocent's coffee and says, 'give me a dollar or I'm keeping this'. Innocent prefers not to lose her coffee and hands over a dollar.

Innocent's consent is clearly invalid; this is not a legitimate transfer of property and her dollar should be returned to her. That said, if we set aside the rights violation, Innocent's options are surely both acceptable; whether she loses her coffee or a dollar, she is at worst mildly inconvenienced. The common feature that invalidates consent in *ROBBERY* and *COFFEE* is thus not the mere absence of acceptable alternatives, but the presence of a specific sort of rights violation: Victim and Innocent each have a right to a third alternative they are unjustly denied, and indeed, they are unjustly denied by the very party soliciting their consent. This hypothesis also explains why Seller's consent is valid in *RING*; it would be very nice for Buyer to simply give Seller the money, but Seller has no right that Buyer do this.

This is, at any rate, how I analyze these cases. That said, sorting out exactly when and why various features of scenarios in which somebody lacks acceptable alternatives invalidate their consent is a task I cannot properly undertake here; this is a lengthy dialectic that invites iterative moves and counter moves. Fortunately, we do not need a perfectly general solution to this (if there even is one to be found). This much is beyond dispute: lacking acceptable alternatives does not in general invalidate somebody's consent. So, what is it about the specific lack of acceptable alternatives to data practices that is thought so problematic?

I believe there is a straightforward answer here. Assume for sake of argument that forgoing certain technologies really is an 'unacceptable alternative', in whatever precise sense you like. Assume also, as is currently true, that people must consent to some agreement to access these technologies. Consent to this agreement is not necessarily invalid. For example, if the agreement simply required users to pay some modest and reasonable fee to the service provider, most users would surely have no grounds for complaint. The mere fact that users lack acceptable alternatives is thus not itself a problem here. But of course, as matters stand, this is not the agreement; users are instead required to accept invasive data practices. Why is this a problem when paying a modest and reasonable fee would not be? I propose that the most plausible answer is that this specific arrangement violates people's

rights, not because they lack acceptable alternatives in the abstract, but because they lack a specific alternative that, by right, should be available to them: one in which they both have access to important technologies and maintain reasonable privacy. Call this ‘the access right’. I believe the best version of the ‘no acceptable alternatives’ argument charges (certain) data practices with violating the access right.

Do members of our society have such a right? I find it quite plausible that they do, as long as we are sufficiently flexible about what counts as ‘reasonable privacy’. But therein lies the rub. It has long been recognized in the literature on the right to privacy that people do not have a right to absolute privacy; they have a right to reasonable privacy, and what counts as reasonable privacy is determined by balancing people’s privacy interests against those various interests of others with which they conflict.³² But this means we have essentially reached the same point here that we did when considering information. The scope and limits of the access right, assuming it is a genuine right at all, depend on the independent justifications that can be offered for and against various data practices. Insofar as these practices are independently justified, they fall outside the scope of the access right, and the access right thus cannot be cited as invalidating user consent to them.

I should mention briefly that there is another way of reading the ‘no choice’ argument according to which the claim is not that consent is invalidated, but simply that consent is rendered pointless. For reasons discussed in the last section, I believe this argument is also unsuccessful; insofar as users can bring meaningful preferences to bear on consent decisions about data practices, and insofar as they have options that variously (even if imperfectly) conform to these preferences, their consent is not pointless. Note also that these conditions need not be satisfied in every case; a policy for soliciting user consent is (minimally) not pointless if users can meaningfully exercise their discretion in some cases, even if they cannot do so for all of them.

4. Externalities

In the space that remains, I wish to briefly say something about the second objection to notice-and-consent, which targets its presupposition that data restrictions are imposed almost exclusively for the protection of individual users. Against this, as Susser says, many critics have argued that ‘... what is at issue in [data practices] is often a social rather than individual interest. In such cases individuals should not be deciding about it in the first place’.³³

The point that data practices can affect third-party interests, including perhaps broad social interests, is well taken. Here are three examples. First, at an individual level, the collection and use of one person’s data can facilitate inferences about nonconsenting third parties, thus potentially undermining their privacy. For example, if enough people voluntarily disclose their sexual orientations, this might facilitate reliable inferences about the sexual orientations of people who did not (and do not want to) reveal this. Second, this time at a broader social level, there may be important values – like autonomy and freedom of speech – that mass surveillance is inimical to. For example, Richards (2013) argues that a society in which everybody’s research habits (e.g. search histories) are constantly monitored may be one in which people are driven to intellectual conformity. Finally, mass surveillance may lead to dangerous amalgamations of power in various entities, private and

public. For example, Richards also argues that mass surveillance can facilitate blackmail, discrimination, manipulation, and the selective enforcement of laws.³⁴

I do not dispute the importance of these arguments, although I worry about some of the details. I do, however, dispute the claim that individuals should automatically be disabled from making some decision merely because it affects third-party interests (whether of other individuals or of society generally). Our decisions, involving consent or otherwise, very frequently affect third-party interests, e.g. when I host a dinner party and thus subject my neighbors to increased noise. Sometimes this is legitimate and sometimes it is not, but the mere observation that an interest has been set back does not by itself make the case. This is true even for broad social interests, especially when these interests are affected not by unilateral decree, but in an emergent way from the decisions of various individuals. For example, the ubiquity of cell phones is transforming society in many ways, some big, some small, and not all of them good (consider e.g. how the persistent availability of cell phones transforms work culture). But this does not mean that people should be prohibited from purchasing or using cell phones. Society is comprised of its members, and their decisions can have cumulative implications for social values without any impropriety.

It might be helpful to distinguish two possible justifications for consent here: first, an individual's consent is solicited to legitimate the effects something has on third-party interests; and second, an individual's consent is solicited to legitimate something, and that thing affects third-party interests.

The second style of justification is extremely common and not at all improper. To apply this to the present context, assume certain data practices conflict with certain social values, e.g. because they facilitate reliable inferences about the sensitive information of nonconsenting third parties. This is only part of the justificatory story; we must also consider any values these data practices facilitate, both to the individuals who participate, and (potentially) to society generally. If these are significant enough – e.g. if they make important technologies significantly better – then it may well be worth permitting these data practices, costs to other values notwithstanding. Assuming they should be permitted, the only question that remains is whose consent, if anybody's, should be required to legitimate them. If these practices rely on extensive data collection, there is a strong case to be made, on grounds of both legitimacy and desirability, that the people whose data is employed should at least be included here. There is thus nothing absurd in requiring individual consent for something that affects social values, at least if the purpose of consent is not to justify these effects but is instead to legitimate something that has them as a consequence.

Note also that the first style of justification may sometimes be valid as well, precisely for broad social values, especially those pertaining to privacy. To focus on the matter at hand, many of our privacy norms are essentially a series of informally negotiated boundaries that demarcate certain conventional zones of privacy.³⁵ I do not mean to suggest that privacy is arbitrary, or valuable only in some sort of relativistic sense. But I do mean to suggest that the precise contours of our privacy norms are to some extent conventional, and insofar as they are, it is legitimate for society, especially as a collective enterprise, to redraw them. Individuals systematically choosing to waive entitlements that fall within this zone may be one legitimate mechanism for doing this. I do not want to overstate the case here, as this process can go awry for various reasons (especially when people have weak information) and its powers of legitimation are not unlimited. But it is another respect in which broad

social values are not in principle beyond the legitimate purview of individual consent decisions.

5. Conclusion

I conclude that user consent to data practices is in principle considerably more robust, and more important, than the standard criticisms of notice-and-consent suggest. This does not mean that notice-and-consent makes for good data policy; data subjects are vulnerable to exploitation, and there are good reasons to suspect they have been (and will continue to be) taken advantage of absent substantive legislation that protects them (although there are also some reasons for optimism). What it does mean is that user consent can, in the right circumstances, legitimate data practices which would otherwise be illegitimate – even complicated ones which are extremely costly to avoid. It also means that user consent must be solicited if such practices are to be deployed legitimately (and to reiterate, there are sometimes good reasons that favor doing so). It is thus important that our data policies continue to emphasize user consent, although exactly what this should look like (e.g. which auxiliary protections should also be implemented and how significantly these should deviate from the current system) is not a matter I have explored here.

Kevin Mills, Department of Linguistics and Philosophy, Massachusetts Institute of Technology, Cambridge, MA, USA. kevinaimills@gmail.com

Acknowledgements

I have received valuable feedback on this article, or the ideas it resulted from, from more people than I can remember. To focus on those I can remember (with my sincere apologies to those I have forgotten), I wish to thank: Jinyoung Baek, John Basl, Eugene Chao, Daniel Dake, Will Fleisher, Kathy He, Anneliese Mills, Sean Murphy, Anton Rubisov, Ronald Sandler, Rory Smead, Christo Wilson, and Allen Wood. My thanks also to two anonymous reviewers for the *Journal of Applied Philosophy* for some extremely helpful feedback; the article is much stronger because of it.

NOTES

- 1 Susser, Daniel. 2019. "Notice After Notice-and-Consent: Why Privacy Disclosures Are Valuable Even If Consent Frameworks Are Not." *Journal of Information Policy* 9: 37–8.
- 2 Ibid., p. 48.
- 3 See e.g. Barocas, Solon, and Helen Nissenbaum. 2014. "Big Data's End Run around Anonymity and Consent." In *Privacy, Big Data, and the Public Good: Frameworks for Engagement* (Vol. 1), edited by Julia Lane, Victoria Stodden, Stefan Bender, and Helen Nissenbaum, 44–75. Cambridge: Cambridge University Press. <https://doi.org/10.1017/CBO9781107590205.004>. For a more cautiously optimistic take, see Solove, Daniel. 2013. "Privacy Self-Management and the Consent Dilemma." *Harvard Law Review* 126: 1880–903.
- 4 Susser op. cit., p. 47.
- 5 My thanks to an anonymous reviewer for pressing me on this.
- 6 Note that normative powers are themselves normative entitlements, and people sometimes have normative powers to alter other normative powers. For example, in selling some property, one transfers the normative

- powers associated with property ownership (e.g. the power to sell the property) to another person. The details will not concern us here.
- 7 The claim that rights and duties always correlate with one another is known as the 'correlativity thesis'. In an unqualified form, this thesis is contentious, but I do not assume it in an unqualified form here; I assume only that the duties one party owes to another correlate with (claim-)rights held by the latter against the former (I do not assume that all rights or duties always so correlate). For a defense of unrestricted correlativity, see especially Kramer op. cit. For objections, see e.g. Raz op. cit., pp. 165–93, and Sreenivasan, Gopal. 2010. "Duties and Their Direction." *Ethics* 120(3): 465–94.
 - 8 See e.g. Susser op. cit. Cf. Nissenbaum, Helen. 2011. "A Contextual Approach to Privacy Online." *Daedalus* 140(4): 32–48.
 - 9 Nissenbaum op. cit., p. 36.
 - 10 Solove op. cit., pp. 1888–9.
 - 11 Barocas and Nissenbaum op. cit., pp. 59–61.
 - 12 For example, for a discussion of just how much can be learned about people from their Facebook 'likes', see Youyou, Wu, Michal Kosinski, and David Stillwell. 2015. "Computer-Based Personality Judgments Are More Accurate than Those Made by Humans." *Proceedings of the National Academy of Sciences* 112(4): 1036–40.
 - 13 Barocas and Nissenbaum op. cit., pp. 59–61. Cf., for technical details, Crain, Matthew. 2018. "The Limits of Transparency: Data Brokers and Commodification." *New Media & Society* 20(1): 88–104.
 - 14 Solove op. cit., pp. 1889–91.
 - 15 See especially Miller, Franklin G., and Alan Wertheimer. 2010. "Preface to a Theory of Consent Transactions: Beyond Valid Consent." In *The Ethics of Consent: Theory and Practice*, edited by Franklin G. Miller and Alan Wertheimer, 79–105. New York: Oxford University Press. See also Wertheimer, Alan. 1996. *Exploitation*. Princeton, NJ: Princeton University Press, 37–76.
 - 16 See e.g. Faden, Ruth R., and Tom L. Beauchamp. 1986. *A History and Theory of Informed Consent*. New York: Oxford University Press, 235–73, 337–82. Notably, they do not explicitly require knowledge of alternatives, although many of the historical proposals they discuss do. Cf. Dougherty, Tom. 2020. "Informed Consent, Disclosure, and Understanding." *Philosophy and Public Affairs* 48(2): 119–50.
 - 17 See especially Faden and Beauchamp op. cit., pp. 114–50.
 - 18 The benchmark claim is from Faden and Beauchamp op. cit., p. 284; the quote is from p. 248.
 - 19 Barocas and Nissenbaum op. cit., pp. 57–8; note that I combined separate passages here, although I believe faithfully.
 - 20 One might object that Culpable's consent is actually invalid but is legitimately taken as valid by Landlord, given his (justifiable) ignorance of this. I think this misdescribes the situation. The function of Culpable's consent is not (primarily) to legitimize otherwise impermissible actions by Landlord; it is to bind Culpable to various financial obligations. Insofar as Culpable really is bound by these obligations, what Landlord may legitimately do is somewhat beside the point, and I find it clearer to refer to her consent as 'valid' since it has transformed normative relations in precisely the ways it purported to. Even still, I am inclined to agree that Culpable's consent is defective in an important respect, as long as we recognize that this defect does not undermine its transformative power; the central argument I go on to develop is that given our unavoidable cognitive limitations, we must sometimes make peace with consent that is defective like this, in ways that depend heavily on the particularities of the case at issue. My thanks to an anonymous reviewer for pressing me to clarify this.
 - 21 See especially Nissenbaum op. cit.; Solove op. cit.
 - 22 This example is inspired by similar examples from Faden and Beauchamp op. cit., p. 254.
 - 23 See especially Miller and Wertheimer op. cit.
 - 24 See Dougherty op. cit., pp. 132–3. My thanks to an anonymous reviewer for pointing out that Dougherty defends a more qualified view on this matter. Dougherty, Tom. 2021. *The Scope of Consent*. Oxford: Oxford University Press. I hedge my own formulation with 'typically' because I see nothing absurd in the proposal that people might reasonably be granted powers that can be unintentionally exercised; these would, however, presumably be unusual cases.
 - 25 This section draws heavily from Wertheimer op. cit., pp. 37–76.
 - 26 See e.g. Scanlon, Thomas. 2000. *What We Owe to Each Other*. Cambridge, MA: Belknap Press of Harvard University Press, 251–6, for some alternatives.
 - 27 See e.g. Verheyden, Tim, Denny Baert, Lente Van Hee, and Ruben Van Den Heuvel. 2019. "Google Employees Are Eavesdropping, Even in Your Living Room, VRT NWS Has Discovered." VRT NWS, July 10, 2019. <https://www.vrt.be/vrtnws/en/2019/07/10/google-employees-are-eavesdropping-even-in-flemish-living-rooms/>

- 28 See e.g. Faden and Beauchamp op. cit., pp. 344–6; Wertheimer op. cit., pp. 52–3; Hawkins, Jennifer S., and Ezekiel J. Emanuel. 2008. *Exploitation and Developing Countries: The Ethics of Clinical Research*. Princeton, NJ: Princeton University Press, p. 25.
- 29 This example is essentially from Wertheimer op. cit., pp. 52–3.
- 30 I want to thank Anneliese Mills for pressing me on this.
- 31 Wertheimer op. cit., pp. 52–3.
- 32 See e.g. Marmor, Andrei. 2015. “What Is the Right to Privacy?” *Philosophy and Public Affairs* 43(1): 3–26; Parent, W. A. 1983. “Privacy, Morality, and the Law.” *Philosophy and Public Affairs* 12(4): 269–88; and Gavison, Ruth. 1980. “Privacy and the Limits of Law.” *Yale Law Journal* 89(3): 421–71. The latter two do not actually use the expression ‘reasonable privacy’, but what they do say seems clearly to express the same basic idea.
- 33 Susser op. cit., pp. 37–8.
- 34 See Richards. 2013. “The Dangers of Surveillance.” *Harvard Law Review* 126: 1934–65.
- 35 See Scanlon, T. M. 1975. “Thomson on Privacy.” *Philosophy and Public Affairs* 4(4): 315–22.